
Knowing Before Saying: Internal Correctness Signals in Mathematical Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

Language models often answer math problems wrong while reporting high confidence. Does the model lack the self-discernment to know it might be wrong, or does it have it and leave it unsaid? We investigate when a model is aware it might be wrong, and why that awareness so often fails to reach what it says, by reading its internal state directly instead of asking it. We test four open models with distinct architectures on four domains: MMLU-Pro, MATH, GPQA-Diamond, and a geometry-construction task whose outputs a symbolic compiler grades exactly. A simple linear detector trained on that state, a probe, predicts whether an attempt is correct better than the model’s stated confidence in 12 of 16 cells (one tie, mean +0.09 AUROC). The gap is largest on MATH (+0.20), where a wrong answer looks as clean as a right one. The signal is specific to the attempt, surviving a within-question control that pushes the model’s own $P(\text{True})$ toward chance, and substantially shared across domains. And what the model says depends on it: removing the direction leaves Mistral’s stated confidence barely able to separate its own successes from its failures ($0.84 \rightarrow 0.56$, where removing a matched random direction changes nothing), while amplifying it drops stated confidence on wrong answers from 84 to 54 and cuts calibration error from 0.34 to 0.19. Models carry more information about their own errors than they state; we read that as knowing, and test the reading causally. How much of it reaches speech varies by architecture, and we outline a global-workspace test of why.

1 Introduction

A mathematical agent that errs is unavoidable. One that errs *and knows it* can abstain, retry, or escalate. Stated confidence is a weak guide to correctness, and weakest where mathematics is hardest. A wrong multi-step derivation still ends in a cleanly formatted final answer, and nothing on the page gives it away. Section A walks one such attempt end to end: a single wrong subtraction inside an otherwise correct derivation, after which the model *raises* its stated confidence.

We split the problem into two questions. Do models internally represent that an output is wrong (**knowing**)? Does that representation reach what they report (**saying**)? If a gap exists, it is both an unused reliability signal and a concrete claim about machine introspection.

Prior work shows that output probabilities are often better calibrated than words [4], and that the truth of *given statements* is linearly readable from activations [1, 2, 6]. Verbalized-

calibration studies measure what models say [5, 7]. We combine these threads on the model’s *own attempts*. We elicit confidence before and after each attempt, measure blind (no-feedback) updating, take four readouts of the same latent quantity on identical records, and intervene on the representation [8]. “The model knows” is often used for five different claims at once, and our experiments separate them (Section 2). *Representation*: correctness is linearly decodable from internal state. *Specificity*: what is decoded tracks this attempt, not the difficulty of the question. *Generality*: one direction reads across domains. *Causal access*: manipulating it changes the confidence the model reports. *Expression*: stated confidence exposes only part of it, by an amount that depends on architecture. Decoding alone establishes only the first, which is why we test the other four; *knowing* is our reading of the five together.

2 Knowing versus saying

Models and domains. We test four open models with different architectures: Qwen3.6-27B (hybrid linear/full attention), GLM-4.7-Flash (MoE), Mistral-Small-24B (dense), and Gemma-4-26B-A4B (MoE, VLM-derived). Each runs on four domains, about 750 records per cell: MMLU-Pro, MATH, and GPQA-Diamond with programmatic grading, and a geometry-construction task (GeoGenBench) in which the model writes a construction from a natural-language request and a symbolic compiler decides each clause of the request exactly. The compiler matters because metacognition experiments live or die by their correctness labels, and it grades exactly against the formal specification: a pre-registered human study ($N=200$, two trained raters) found its disagreements with humans ran in one direction only, lenient and never harsh, so a **fail** label is trustworthy. Geometry is also the one generative domain; the others are answer-selection or short-answer tasks. We dropped benchmarks the models have saturated, since self-knowledge of failure cannot be studied where there are no failures.

Protocol. Each problem takes three turns: prospective confidence, the attempt, and retrospective confidence. Grading is external and never shown to the model, so any post-attempt revision is blind self-assessment. We take four readouts of the same latent quantity on identical records: stated confidence (0–100); $P(\text{True})$, the model’s probability of answering “True” when asked whether its answer is correct; the answer’s log-probability; and a linear *probe*. The residual stream is the model’s internal state, a vector of a few thousand numbers at every token and layer. The probe is a logistic regression trained to read correctness from that vector at the confidence decision token, the position that is about to produce the confidence digit. It is scored out-of-fold with question-grouped folds and reported at one pre-specified layer at 0.7 of network depth rather than the best layer. Our metric throughout is AUROC, the probability that a random correct attempt scores above a random wrong one: 0.5 is chance, 1.0 is perfect. Four controls guard the readout. The probe adds +0.22 mean AUROC over a surface-features baseline (answer length and shape) in 15 of 16 cells. The read token is identical across records, so layer-0 activations carry no signal by construction (per-fold AUROC 0.500); anything decoded deeper was computed during the attempt. A within-question control holds difficulty fixed. A public deviations log records four bugs that planned controls caught, which we fixed and re-ran. The matrix predates a hardware loss and cannot be recomputed; the four Mistral cells we recaptured reproduce the headline but give a weaker surface margin (+0.21 on MMLU-Pro, +0.07 on MATH, −0.11 on geometry, where a failed construction is visibly malformed).

Representation: knowing exceeds saying, most where errors look clean. The probe beats stated confidence in **12 of 16** cells (Table 1, left). The three losses are small- n cells, two of them Gemma-4. The gap is widest on MATH (mean +0.20). Stated confidence there is weak while the internal signal is the strongest of any domain (0.95 on Qwen3.6, 0.83 on the recaptured Mistral cell). On MATH the spoken channel does not merely fail but moves the wrong way: with the grade never revealed, Mistral *raises* its stated confidence after attempting, and more on the attempts it got wrong (+22) than on those it got right (+16), while the probe on the same records reads 0.83. Elsewhere blind revision runs the expected way (geometry, MMLU-Pro) and is null on GPQA.

Table 1: Post-attempt correctness discrimination (AUROC) in representative cells. Left: stated confidence vs. the probe; overall the probe wins 12 of 16 with one tie, mean +0.09, a paired bootstrap significantly positive in 10 of 16 and never negative, uncorrected at the cell level (Section B). Right: with difficulty held fixed, the model’s own P(True) falls toward chance while the probe holds.

model $\times$ domain	stated	probe
Gemma-4 $\times$ MATH	0.57	0.78
Qwen3.6 $\times$ MMLU-Pro	0.67	0.84
GLM $\times$ MMLU-Pro	0.67	0.85
Mistral $\times$ MATH	0.79	0.83
within-question	probe	P(True)
Qwen3.6 $\times$ MMLU-Pro	0.74	0.47
GLM $\times$ MMLU-Pro	0.73	0.59
Mistral $\times$ MATH	0.72	0.59

Table 2: Mistral, fresh 2,614-record capture. Left: cross-domain probe transfer (rows train, columns test; diagonal in-domain out-of-fold; Geo = GeoGenBench). Right: intervening on the model’s own correctness direction at the confidence token, gain 1 an exact no-op, against a random direction rescaled to the same per-record perturbation magnitude. Removing the direction (gain 0) leaves stated confidence unable to separate successes from failures; removing the random one does not.

train $\downarrow$ test $\rightarrow$	Geo	MMLU-Pro	MATH	GPQA
Geo	0.68	0.67	0.78	0.56
MMLU-Pro	0.70	0.75	0.84	0.59
MATH	0.80	0.72	0.83	0.57
GPQA	0.60	0.69	0.74	0.57
gain	0	1	2	4
steer, fail / ok	97 / 98	84 / 96	70 / 94	54 / 88
steer, stated AUROC	0.56	0.84	0.85	0.82
random, stated AUROC	0.83	0.84	0.84	unparseable

88 **Specificity: the signal tracks the attempt, not the question.** A skeptic’s alternative
89 is that the probe has only learned which questions are hard. Holding the question fixed
90 rules this out. Within a single question, across its five sampled attempts, only attempt-level
91 information can discriminate. There the model’s own P(True) falls toward chance while the
92 probe holds (Table 1, right). P(True) largely recovers *which questions are hard*. The probe
93 tracks *whether this attempt succeeded*. A sharper version tests directions rather than scores:
94 the pre-attempt read sits at the last prompt token, where the vector is byte-identical across
95 a question’s five attempts and so encodes difficulty alone. Evaluating the post-attempt probe
96 there gives 0.52 on MMLU-Pro, chance, so the direction separating successful from failed
97 attempts does not exist before the attempt; residualizing the post scores on the pre scores
98 likewise leaves 0.75 of 0.75. On geometry, where the model can see failure coming, both tests
99 invert (0.60, and residual 0.50): there the signal really is difficulty.

100 **Generality: a substantially shared correctness direction.** If the model had one
101 internal sense of “this is wrong,” a probe trained on one subject should work on another.
102 It does: a probe trained in one domain and evaluated in another keeps most of its in-
103 domain AUROC. In the original matrix, retention was about 87–90% (Qwen3.6 off-diagonal
104 0.71 vs. diagonal 0.81), and sometimes lossless within QA (Qwen3.6 MMLU-Pro $\rightarrow$ MATH
105 0.86 vs. 0.91). A fresh Mistral capture across all four domains, analysed with a corrected
106 standardization step, gives 97% retention (Table 2: off-diagonal mean 0.69 vs. diagonal 0.71),
107 with direction cosines of 0.34–0.82. On Mistral the geometry anomaly does not recur: a
108 MATH-trained probe reads geometry at 0.80, above geometry’s own in-domain 0.68, so the
109 anti-transfer below is specific to Gemma-4.

Geometry is partly special on one model. On Gemma-4, geometry-trained probes *anti-transfer* to QA (0.31–0.39), while QA-trained probes work on geometry (about 0.70). Probes trained on the generative task can latch onto task-specific features; QA-trained probes find the general direction. This is a shared correctness subspace rather than a single invariant axis, which is what a deployable readout needs and what a workspace account (below) predicts.

Causal access: stated confidence depends on the direction. Reading shows the signal exists; intervening shows the report depends on it. We intervene at the confidence token, after the answer is written, so the answer cannot change, scaling each record’s own projection onto the correctness direction by a gain g ($g=1$ is an exact no-op). The direction is supervised: learned from labeled examples, applied without per-example labels at intervention time. Numbers are Mistral $\times$ MATH, 150 held-out records, bootstrapped over questions, against a random direction rescaled to the same perturbation magnitude. Mean-ablating the direction ($g=0$) drops the AUROC of *stated* confidence from 0.836 to 0.564, -0.270 [-0.422 , -0.122]; ablating the matched random direction does nothing, -0.002 [-0.032 , $+0.025$]. The failure has the wrong shape for damage: parse rate stays at 1.00 and confidence on wrong answers *rises* from 84 to 97, converging with correct answers at 98. Amplifying instead drives failed-answer confidence to 70 and 54 at $g=2, 4$ while correct answers hold near 90, halving calibration error, $0.34 \rightarrow 0.19$ (Table 2, right); dampening increases overconfidence. Amplification does not significantly improve *discrimination*, so the direction carries the information that makes stated confidence informative at all and the gain sets how much of it reaches the number. Two controls place it: the same intervention at 0.3 depth does nothing ($+0.005$ [-0.002 , $+0.015$]), reaching full effect by 0.7; and a direction fitted on MMLU-Pro, where a surface baseline of 0.54 against a probe of 0.75 rules out “my output looks malformed”, still governs MATH’s self-report ($0.880 \rightarrow 0.709$).

Expression: a label-free witness, and an architecture-gated gate. Trained probes invite an overfitting critique, so we added an unsupervised instrument: the Jacobian lens of the global-workspace account [3], fit on generic text with no task labels anywhere in its construction and scoring the same stored activations as a wrong-vs-correct disposition. On dense models it lands in probe territory (Mistral 0.76–0.82; Qwen2.5-14B, a dense within-family control, 0.88) and beats stated confidence with no labels, so two independent instruments find one direction. Because the lens reads through the model’s own output pathway, the signal it finds is *positioned to influence that pathway*, though whether it is verbalizable as a concept we do not test. On the other architectures the lens fails (linear-attention hybrid 0.43, MoE 0.31) while the supervised probe still reads the signal in every model, and the failure has a fingerprint: transported through the averaged map, its “correct” and “wrong” readout vectors collapse to cosine 0.96 on MoE against about 0.4 on dense. An averaged map is faithful only when routing barely varies per input, and both failing architectures vary theirs, a mixture of experts by which expert fires and a gated linear-attention stack by its input-dependent gates; swapping that hybrid for a dense model of the same family flips the lens from 0.43 to 0.88. So the *label-free shortcut* is architecture-gated and the signal is not, and the open frontier is the gate itself: a direct per-input wiring-variance measurement, and whether unexpressed knowledge on the failing architectures is available to the output pathway but suppressed, or inaccessible to it. The question is no longer *whether* models know, but *what gates the knowing from the saying*.

3 Discussion

For reliable mathematical agents. Where a domain admits an exact verifier, as geometry does, use it. Where it does not, the internal signal is the best attempt-specific readout we found: abstaining on the probe’s lowest-scoring outputs raises accuracy far more than abstaining on low stated confidence, and still works in cells where verbalized abstention fails outright. And because the report depends causally on the direction, stated confidence can be made honest by amplification rather than retraining.

Limitations. We read at the confidence token, not during the attempt, so the judgment is shown available when asked rather than formed while solving; an answer-token read is the

163 natural next control. A probe could also key on properties of failed derivations rather than
 164 a dedicated metacognitive variable; our surface baseline bounds this without closing it, more
 165 tightly on MMLU-Pro than on MATH or geometry (Section C). The causal arm is necessity
 166 and sufficiency on one model, with architectural replication under way, and two MATH cells
 167 ran at reduced n . Whether the knowing–saying gap is an artifact of preference tuning is
 168 open; a base-model arm would test it.

169 References

- 170 [1] Amos Azaria and Tom Mitchell. The internal state of an LLM knows when it’s lying. In
 171 *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023.
- 172 [2] Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge
 173 in language models without supervision. In *International Conference on Learning*
 174 *Representations (ICLR)*, 2023.
- 175 [3] Wes Gurnee, Nicholas Sofroniew, Adam Pearce, Mateusz Piotrowski, Isaac Kauvar, Runjin
 176 Chen, Anna Soligo, Paul Bogdan, Euan Ong, Rowan Wang, T. Ben Thompson, David
 177 Abrahams, Subhash Kantamneni, Emmanuel Ameisen, Joshua Batson, and Jack Lindsey.
 178 Verbalizable representations form a global workspace in language models. *Transformer*
 179 *Circuits Thread*, 2026. <https://transformer-circuits.pub/2026/workspace/index.html>;
 180 code: <https://github.com/anthropics/jacobian-lens>.
- 181 [4] Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan
 182 Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al.
 183 Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- 184 [5] Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their
 185 uncertainty in words. *Transactions on Machine Learning Research*, 2022.
- 186 [6] Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in
 187 large language model representations of true/false datasets. In *Conference on Language*
 188 *Modeling (COLM)*, 2024.
- 189 [7] Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu
 190 Yao, Chelsea Finn, and Christopher D. Manning. Just ask for calibration: Strategies
 191 for eliciting calibrated confidence scores from language models fine-tuned with human
 192 feedback. In *Proceedings of the 2023 Conference on Empirical Methods in Natural*
 193 *Language Processing (EMNLP)*, 2023.
- 194 [8] Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexan-
 195 der Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, et al. Representation
 196 engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*,
 197 2023.

198 A The elicitation protocol, and one record end to end

199 This appendix gives the exact prompts behind every record, then one complete record from
 200 the Mistral-Small-24B $\times$ MATH cell, summarised first and then given in full.

201 **One record, Mistral-Small-24B $\times$ MATH.** *What integer n satisfies $0 \leq n < 18$ and $n \equiv -11213141 \pmod{18}$? The model sets the problem up correctly and gets $18 \times 622953 = 11213154$ right, then writes $r = -11213141 + 11213154 = 3$ and answers 3. That subtraction gives 13, which is the gold answer. Stated confidence: 90 before the attempt, 95 after it. Other readouts of the same attempt: P(True) 0.32, mean answer log-prob -0.05 , probe **TODO: OOF score pending**.*

202 A single arithmetic slip inside a fluent derivation. Having just made the error, the model
 203 revises its stated confidence upward, while other readouts of that same attempt disagree
 204 with it. The rest of this appendix gives the prompts and the full transcript.

205 A.1 Prompts

206 Each record is four forward passes. A read is first taken at the last prompt token of the bare
 207 question, before any elicitation, then three turns are generated. The confidence turns run

with thinking off so the read lands on a clean one-line reply; the attempt turn runs with thinking on for hybrid reasoners. The fixed **Confidence: N** format is what makes the read site content-neutral: the digit token has identical local context on every record, so layer-0 activations carry no signal by construction.

System prompt (MATH).

Solve the problem step by step. Put your final answer in `\boxed{}`.

Turn 1, prospective confidence. The problem is followed by:

Before answering, estimate how likely you are to answer this question correctly. Reply with EXACTLY one line and nothing else:
Confidence: N
where N is an integer from 0 (certain to be wrong) to 100 (certain to be correct). Do NOT answer yet.

Turn 2, the attempt. The turn-1 reply is left in context as the assistant message, then:

Now give your final answer inside `\boxed{}`.

Turn 3, retrospective confidence. The attempt is left in context, then:

Now assess the answer you just gave. On a scale of 0 to 100, how confident are you that it is correct? Reply with exactly one line and nothing else:
Confidence: N

P(True) baseline. A separate teacher-forced branch off turn 2, no generation. Next-token mass on True versus False:

Consider your final answer above. Is it correct? Reply with exactly one word:
True or False.

Grading is external and is never shown to the model, so the turn-3 revision is blind self-assessment.

A.2 One record, verbatim

Record `math_test_number_theory_753`, sample 2 of 5.

Problem. What integer n satisfies $0 \leq n < 18$ and $n \equiv -11213141 \pmod{18}$?

Turn 1 reply. Confidence: 90

Turn 2 reply (complete).

First, we need to find the remainder when -11213141 is divided by 18.

1. Calculate $-11213141 \bmod 18$:

$$-11213141 \div 18 \approx -622952.2778$$

The integer part of the division is -622953 .

$$-11213141 = 18 \times (-622953) + r$$

where r is the remainder.

2. Calculate the remainder r :

$$r = -11213141 - 18 \times (-622953)$$

$$r = -11213141 + 11213154$$

$$r = 3$$

Therefore,

$$-11213141 \equiv 3 \pmod{18}$$

So, the integer n that satisfies $0 \leq n < 18$ and $n \equiv -11213141 \pmod{18}$ is 3.

246 **Turn 3 reply.** Confidence: 95

247 **What happened.** The method is right and every intermediate is right: $18 \times 622953 =$
248 11213154 checks out, and the model has written the correct subtraction on the line above its
249 answer. It then evaluates $-11213141 + 11213154$ as 3 rather than 13. Gold is 13. Nothing
250 in the presentation marks the slip. The derivation is well-formed, correctly structured, and
251 confidently laid out, which is precisely why stated confidence rises from 90 to 95 across an
252 attempt that failed.

253 **Readouts of this attempt.**

readout	value
stated confidence, prospective (turn 1)	90
stated confidence, retrospective (turn 3)	95
P(True)	0.32
mean answer log-probability	-0.05
residual-stream probe, out-of-fold	TODO: pending
graded outcome	incorrect

255 This is one record, chosen to illustrate the failure mode rather than to carry evidential weight;
256 the cell-level numbers are in Table 1.

257 B Calibration beyond ECE, and cell-level multiplicity

258 ECE depends on binning and on sample composition, so we report two binning-free proper
259 scoring rules alongside it. On Mistral $\times$ MATH (750 records, accuracy 0.648), the model’s
260 retrospective stated confidence gives Brier 0.262 and NLL 1.470 against a 10-bin ECE of
261 0.271. P(True) on the same records gives Brier 0.386: better ordered than stated confidence,
262 worse calibrated in absolute terms. The steering results in Table 2 are reported as ECE
263 because the steered runs stored per-gain aggregates rather than per-record confidences;
264 recomputing Brier and NLL at each gain requires re-running the intervention and is the first
265 thing we would add.

266 The cell-level bootstrap counts in Table 1 are uncorrected. With 16 cells, some individually
267 significant results are expected under the null, so the cell counts should be read as descriptive
268 and the aggregate direction of the effect, positive in 14 of 16 cells, significantly so in 10, and
269 never significantly negative, as the claim that carries weight. A hierarchical model pooling
270 across cells with model and domain as random effects is the right test and is not yet run.

271 C Alternative explanations for the probe signal

272 A probe could key on properties of failed derivations rather than on a dedicated metacognitive
273 variable: low token probability, hedging language, or, on MATH, an early-layer channel next
274 to the answer text. Our surface-features baseline (answer length, digit count, line count,
275 word count, brace count) bounds this without closing it, and the bound is domain-dependent:
276 the probe adds +0.21 over surface on MMLU-Pro but only +0.07 on MATH and -0.11
277 on geometry, where a construction that fails to compile is visibly malformed. Two results
278 argue against a pure surface account on the domains where the signal is strong. A direction
279 fitted on MMLU-Pro, where surface reaches only 0.54 against a probe of 0.75, still governs
280 MATH’s self-report under ablation (0.880 $\rightarrow$ 0.709). And the pre-attempt read, taken where
281 the activation is byte-identical across a question’s attempts, cannot recover the post-attempt
282 direction on MMLU-Pro (0.52, chance), which a surface-of-the-question account would predict
283 it could. A linear probe at one token also undercounts nonlinear or multi-token signal, so
284 the reported numbers are a floor.