
GeoGenBench: Symbolically Verified Generation of Geometry Diagrams

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Geometry benchmarks ask a model to read a diagram. None asks it to
2 draw one. GEOGENBENCH does. It has 801 prompts, each asking for a
3 plane-geometry construction and each paired with the properties a correct
4 drawing must have. The model writes TikZ. We compile it, recover the
5 coordinates, and check every property exactly. No human grades and no
6 vision model judges. We test six frontier models. Asking for a typed
7 intermediate representation instead of raw TikZ lifts every model, by 2 to 25
8 points, most for the weakest. Under that scaffold the three best models tie
9 within one point at a $7\times$ cost spread. A human check of 200 outputs finds
10 the verifier lenient in one direction only: it never fails a correct drawing.
11 We release the dataset, the templates, the verifier, and the leaderboard.

12 1 Introduction

13 Language models now write worksheets and slides, and in school the picture they most often
14 need to draw is a geometry diagram. Every geometry benchmark tests the other direction:
15 here is a diagram, answer a question. None tests whether a model can draw a correct diagram
16 from a description.

17 Take the prompt “*Draw an acute triangle EFG with angles 60° and 70° at E and F .
18 Construct the altitude from G meeting EF at H .*” A correct drawing has those two angles,
19 H on the segment EF and not beyond it, a right angle at H , and no obtuse angle. Each of
20 those is a check we can run on the coordinates of the model’s drawing. In our pilot, one
21 flagship model put H beyond the segment. Another drew an obtuse triangle. The verifier
22 names the check that failed.

23 Two ideas make this work. The prompt and its checks come from the same template, so the
24 checks can never ask for something the prompt did not. And the drawing is graded from its
25 coordinates, so “right angle” is a calculation, not an opinion about how the picture looks.

26 **Prior work.** Geometry3K [6], GeoQA [3], UniGeo [4], MathVista [7], GeoEval [10], and
27 GeoBench [1] test diagram reading. T2I-CompBench++ [5] and GenExam [9] grade generated
28 images with a model as judge, which can say whether a picture looks like a right triangle but
29 not whether it has a right angle. AlphaGeometry [8] and the Lean corpora machine-check
30 proofs. GEOGENBENCH machine-checks drawings.




Figure 1: Procedural co-generation: one Python template emits both the prompt and the rubric, guaranteeing they stay mutually consistent.

2 The benchmark

Co-generation. One Python template writes both the prompt and its checks (Figure 1). The altitude template returns the sentence the model sees and the list `[right_angle(A, H, B), point_on_segment(H, B, C)]` the grader uses. The two cannot drift apart. A reviewer worried about contamination can re-roll the template with a fresh seed.

Symbolic verification. The model emits TikZ. We compile it twice: once to an image, once to SymPy geometry with coordinates. Fifteen checks such as `right_angle`, `midpoint`, `parallel`, and `point_on_segment` run on those coordinates in microseconds, with a small numeric tolerance. Two weaker checks on labels and marks count only toward a loose pass. Template authors say what must hold, not how to test it.

Scoring. The model sees only the prompt. An output passes if every check passes. It gets a soft pass if nothing fails but a check was skipped. Otherwise it fails, or it never rendered.

Composition. 600 of the 801 prompts come from 30 templates, 200 per tier: T1 is a single shape, T2 adds one construction, T3 combines several. The other 201 were extracted by a language model from a state K–12 curriculum [2]. We report the two sources separately and treat the templated 600 as the anchor.

3 Verifier reliability

Completeness. A check list is complete if every drawing that passes it is the intended construction, up to position and scale. Of the 30 templates, 18 are that tight. The other 12 pin the named relations but leave free parameters, as a teacher’s instructions usually do. We also tried to fool each check with TikZ that passes it while contradicting the prompt. Five checks can be fooled alone. Templates pair checks, so we estimate under 5% of passes are exploitable.

Hardening. Three verifier fixes promoted 4.6% of pilot outputs from soft pass to pass and demoted none. They also halved the apparent gap between the two strategies. Part of that gap had been the verifier’s, not the models’. Every number below is post-fix.

Human check. Two raters, a PhD mathematician and an engineer who has taught mathematics, judged 200 outputs blind to the verifier’s verdict. Verifier and raters agree moderately ($\kappa = 0.455$), below the 0.70 we had pre-registered. The raters agree well with each other, so the gap is the verifier’s. The disagreement runs one way. The raters overturned 18% of the verifier’s passes and none of its fails. Almost all of those sit on T2, where the checks confirm the construction but not the numbers the prompt gave. Leaderboard rates are therefore a little generous, evenly across models. Numeric checks for those templates are the planned fix.

Table 1: Headline leaderboard (600 scenarios $\times$ 3 repeats, post-hardening). Strict = **pass**; loose adds **soft_pass**. * marginally partial coverage; $\dagger$ blocked at the API layer (Section 4).

Model	Strategy	N	Strict	Loose	\$/scen.	Latency
GPT-5.5	structured	2,304*	94.1%	94.1%	\$0.075	33.6 s
Claude Sonnet 4.6	structured	2,861	93.4%	93.4%	\$0.088	19.0 s
Claude Opus 4.7	structured	2,292*	93.3%	93.3%	\$0.532	23.9 s
GPT-5.5	raw_code	1,800	91.9%	95.3%	\$0.058	14.8 s
Gemini 2.5 Pro	structured	1,140*	83.1%	83.1%	\$0.050	75.7 s
Claude Sonnet 4.6	raw_code	1,800	81.2%	95.3%	\$0.044	16.4 s
Claude Haiku 4.5	structured	2,232	75.3%	76.8%	\$0.032	8.0 s
Claude Opus 4.7	raw_code	1,800	71.6%	95.3%	\$0.262	13.8 s
Claude Haiku 4.5	raw_code	2,403	70.5%	76.0%	\$0.019	9.5 s
Gemini 2.5 Pro	raw_code	1,800	58.4%	89.9%	\$0.049	32.9 s
Gemini 2.5 Flash	raw_code	2,403	35.9%	42.2%	\$0.012	22.1 s
Gemini 2.5 Flash	structured	0 $\dagger$	—	—	—	—

4 Experiments

Setup. Six models from three vendors at a $20\times$ price spread: Claude Opus 4.7, Claude Sonnet 4.6, Claude Haiku 4.5, GPT-5.5, Gemini 2.5 Pro, and Gemini 2.5 Flash. Two strategies. `raw_code` hands the model the prompt and a TikZ tutorial and takes what it writes. `structured` asks for typed JSON describing points, segments, circles, and checks, which our code turns into SymPy and then TikZ. The checks run before rendering, so the model gets up to three retries. Every model-by-strategy cell runs on the 600 templated prompts three times.

The scaffold lifts every model. `structured` helps each model, from 2 points for GPT-5.5 to 25 for Gemini 2.5 Pro. The weaker a model is at raw TikZ, the more it gains.

The frontier converges. With the scaffold, GPT-5.5, Claude Sonnet 4.6, and Claude Opus 4.7 finish within one point of each other, at 94%, 93%, and 93%. Claude Opus 4.7 costs seven times what GPT-5.5 does. Without the scaffold the spread runs from 92% down to 36%. T1 is nearly solved. T2 is where the scaffold first matters. T3 is where the frontier models separate. Most wrong outputs are clean diagrams with the wrong geometry, not broken code.

Portability. `structured` needs schema-constrained decoding. Our schema compiles to thousands of states, and one vendor’s API refuses it before the model writes a token. That is a gap in the API, not the model, and a warning for any benchmark built on a large schema.

5 Conclusion

Limitations. GEOGENBENCH tests whether a model can turn a sentence into a consistent plane drawing. It does not test proofs, solids, or teaching. A third of the templates accept any drawing that meets the stated constraints, not one canonical drawing. The curriculum prompts are harder for every model than the templated ones, and the headline numbers are templated numbers.

Takeaway. A geometry drawing can be graded exactly, from its coordinates, with no judge. Graded that way, the scaffold matters more than the model. We release everything needed to run it.

References

- [1] Anonymous. GeoBench: Hierarchical geometric reasoning benchmark, 2026. ICLR 2026 submission, under review.

- 97 [2] Carnegie Learning, Inc. and Texas Education Agency. Bluebonnet learning mathematics:
 98 Teacher’s edition (grades 6–8 and geometry). Curriculum materials adopted under the
 99 Texas State Board of Education Instructional Materials and Technology Allotment,
 100 2024. Accessed via the Texas Education Agency Bluebonnet Learning programme;
 101 <https://learn.texas.gov/>.
- 102 [3] Jiaqi Chen, Jianheng Tang, Jinghui Qin, Xiaodan Liang, Lingbo Liu, Eric P. Xing, and
 103 Liang Lin. GeoQA: A geometric question answering benchmark towards multimodal
 104 numerical reasoning. In *Findings of ACL-IJCNLP*, 2021.
- 105 [4] Jiaqi Chen, Tong Li, Jinghui Qin, Pan Lu, Liang Lin, Chongyu Chen, and Xiaodan
 106 Liang. UniGeo: Unifying geometry logical reasoning via reformulating mathematical
 107 expression. In *Proc. Conf. on Empirical Methods in Natural Language Processing*
 108 (*EMNLP*), 2022.
- 109 [5] Kaiyi Huang, Chengqi Duan, Kaiyue Sun, Enze Xie, Zhenguo Li, and Xihui Liu.
 110 T2I-CompBench++: An enhanced and comprehensive benchmark for compositional
 111 text-to-image generation. *IEEE Trans. Pattern Anal. Mach. Intell. (TPAMI)*, 2025.
- 112 [6] Pan Lu, Ran Gong, Shibiao Jiang, Liang Qiu, Siyuan Huang, Xiaodan Liang, and Song-
 113 Chun Zhu. Inter-GPS: Interpretable geometry problem solving with formal language
 114 and symbolic reasoning. In *Proc. Annual Meeting of the Association for Computational*
 115 *Linguistics (ACL)*, 2021.
- 116 [7] Pan Lu, Hritik Bansal, Tony Xia, Jiacheng Liu, Chunyuan Li, Hannaneh Hajishirzi,
 117 Hao Cheng, Kai-Wei Chang, Michel Galley, and Jianfeng Gao. MathVista: Evaluating
 118 mathematical reasoning of foundation models in visual contexts. In *Proc. Int. Conf. on*
 119 *Learning Representations (ICLR)*, 2024.
- 120 [8] Trieu H. Trinh, Yuhuai Wu, Quoc V. Le, He He, and Thang Luong. Solving olympiad
 121 geometry without human demonstrations. *Nature*, 625:476–482, 2024.
- 122 [9] Zhaokai Wang, Penghao Yin, Xiangyu Zhao, Changyao Tian, Yu Qiao, Wenhai
 123 Wang, Jifeng Dai, and Gen Luo. GenExam: A multidisciplinary text-to-image exam.
 124 *arXiv:2509.14232*, 2025.
- 125 [10] Jiaxin Zhang, Zhongzhi Li, Mingliang Zhang, Fei Yin, Chenglin Liu, and Yashar
 126 Moshfeghi. GeoEval: Benchmark for evaluating LLMs and multi-modal models on
 127 geometry problem-solving. In *Findings of ACL*, 2024.