
Knowing Before Saying: Internal Correctness Signals in Mathematical Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Language models are often confidently wrong. We ask whether a model’s
2 activations carry a better signal of its own errors than its stated confidence
3 does. We test four open models on four tasks: MMLU-Pro, MATH, GPQA-
4 Diamond, and a geometry task graded by a compiler. A linear probe on the
5 hidden state predicts whether an answer is right better than the model’s
6 stated confidence in 12 of 16 cases, by 0.09 AUROC on average. The gap is
7 largest on MATH. The probe is not just detecting hard questions: it still
8 works when the same question is attempted five times. The signal exists
9 before the model is asked how confident it is. On Mistral, the model uses this
10 signal when it reports confidence: erase it from the hidden state and stated
11 confidence no longer separates right answers from wrong ($0.84 \rightarrow 0.56$),
12 while erasing a random direction does nothing. On GLM and Qwen3.6,
13 erasing it changes nothing. In short, these models carry more information
14 about their own errors than they report, and at least one of them uses that
15 information when it reports.

16 1 Introduction

17 A language model that is confidently wrong is a problem. Its stated confidence, which we call
18 its self-report, is a signal for deciding whether to trust its answer. This paper asks whether
19 the model’s activations contain a signal of when it is wrong.

20 Two things could be going on. Upon giving an incorrect answer and saying it has high
21 confidence, either the model genuinely has no idea it might be wrong and is just blindly
22 confident, or it deep down holds the error signal in its activations but does not surface it in
23 its self-report. Hence it is either a capability problem or a reporting problem.

24 In this paper we tell these two apart. We read the model’s hidden state with a linear *probe*
25 and compare what the probe finds to what the model reports. We say a fact is *represented*
26 when a probe can read that fact from the model’s state. We say a fact is *reported* when that
27 fact reaches the confidence number the model writes. We refer to “knowing” as shorthand for
28 a representation that is present, specific to the attempt, and, in at least one model, used by
29 the self-report. We make no claim about awareness. A probe alone could not support one.

30 Prior work reads the truth of a statement from activations [1, 2, 9], and reads a model’s own
31 errors from its hidden state better than its output reveals them [3, 10]. A model’s output
32 probabilities are often better calibrated than its words [5], and asked confidence is known
33 to be overconfident [8, 11, 12]. Probes can latch onto confounds rather than the intended
34 variable [6], which is why Section 3 exists. Intervening on a truth direction changes what a
35 model says [7, 13]. We apply those tools to the model’s own attempts, and then ask whether
36 the self-report depends on the representation.

2 Reading the state

Models and tasks. We test four open models of three architectures: Mistral-Small-24B (dense), Qwen3.6-27B (gated linear attention interleaved with full attention), and GLM-4.7-Flash and Gemma-4-26B (mixture of experts). Each answers 150 questions, five times each, in four domains: MMLU-Pro, MATH, GPQA-Diamond, and GeoGenBench. GeoGenBench is a geometry task. The model writes a construction and a compiler checks it against a formal specification, so its labels need no judge. That is 16 cells of about 750 attempts each.

Prompting structure.

1. The model is shown the question and asked, before answering, how confident it is, from 0 to 100, that it will get it right.
2. It answers.
3. It is asked how confident it is that the answer it just gave is correct.

The probe. At the moment the model is about to write its confidence digit, we record the number and the residual-stream hidden state at that token (2,048 to 5,120 dimensions, depending on the model). The probe is a linear classifier on that vector: features are standardized, reduced to 50 principal components, and fed to an L2-regularized logistic regression that predicts whether the attempt was correct. We fit it with five-fold cross-validation grouped by question, so all five attempts at a question fall in the same fold and the probe is always scored on questions it never saw. The read depth is fixed in advance at 70% of the network (Appendix E). Every readout is scored with AUROC.

The activation state carries more than self-report. The probe beats the model’s self-report in **12 of 16** cells, with one tie. The mean gain is +0.09 AUROC (Appendix D). The gap is largest on MATH, at +0.20. Mistral on MATH is a good example of what this looks like. The model is never told whether it got the answer right. Yet after attempting, it raises its confidence, and it raises it more on the attempts it got wrong (+22 points) than on the ones it got right (+16 points). The probe on the same records reads 0.83. The state tracks correctness better than the self-report does.

Why MATH. On MATH, right and wrong answers have the same shape. Both are a chain of steps ending in a boxed result, so the model cannot tell from the look of its answer whether it went wrong. Appendix A follows one attempt where a single wrong subtraction leads to a wrong answer, and the model then raises its confidence. Shape still carries a little signal. A baseline classifier that sees only surface features of the answer text, such as its length and line count, reaches 0.74, but the probe reaches 0.83. Geometry is the opposite case. A failed construction usually does not compile, so the failure is visible in the answer itself, and there the surface baseline does as well as the probe. The internal signal helps most where the failure is invisible from outside.

Controls. Three things could make a probe look better than it is; the next section tests each. It could be reading how hard the question is rather than the attempt, so we compare attempts within one question. It could be reading the shape of the output, so we compare it with the surface baseline in every cell; on Mistral, adding the probe to the surface features raises AUROC by +0.21 on MMLU-Pro, +0.07 on MATH, +0.05 on GPQA, and −0.11 on geometry. And it could be reading something that was there before the attempt, so we check the input layer, where the token is identical in every record and the probe sits at exactly 0.5. These controls were run on Mistral only, for time. Its four cells reproduce the headline.

3 What the signal is

Difficulty control. A probe that beats self-report could be doing something simpler. It could be learning which questions are hard. Hard questions fail more often, so a difficulty detector would predict correctness without reading the attempt. We run two checks.

Table 1: Mistral. Rows: training domain; columns: test domain; bold: in-domain.

train ↓ test →	GeoGen	MMLU-Pro	MATH	GPQA
GeoGen	0.68	0.67	0.78	0.56
MMLU-Pro	0.70	0.75	0.84	0.59
MATH	0.80	0.72	0.83	0.57
GPQA	0.60	0.69	0.74	0.57

Table 2: Mistral on MATH, 150 held-out records. The correctness direction scaled by a gain at the confidence token, after the answer is written. Bold marks the erased and quadrupled settings; unparseable means the output no longer parsed as a confidence.

	untouched	erased	doubled	quadrupled
gain on the correctness direction	1	0	2	4
mean confidence, wrong / right answers	84 / 96	97 / 98	70 / 94	54 / 88
AUROC of self-report	0.84	0.56	0.85	0.82
same, with a random direction	0.84	0.83	0.84	unparseable

85 First, hold the question fixed. Each question is attempted five times, so within one question,
 86 difficulty cannot vary. The probe still separates the successful attempts from the failed ones
 87 (0.72 to 0.74 in the cells of Table 3). The model’s own $P(\text{True})$, its probability of answering
 88 “True” when asked whether it was right, drops to near chance (0.47 to 0.59).

89 Second, test where the probe’s direction lives. The probe is a direction in the hidden state,
 90 and we can ask when that direction appears. Before the attempt, at the last token of the
 91 question, the model’s state is the same across all five attempts. It can encode difficulty
 92 and nothing else. A probe trained after the attempt and read at that point scores 0.49 on
 93 MMLU-Pro, which is chance, against 0.75 at its own site. The direction that separates good
 94 attempts from bad ones does not exist until the model has tried. Geometry is the exception.
 95 There the pre-attempt state already predicts failure, so that cell is mostly difficulty.

96 **Reading before the question.** A subtler worry. Every read so far came after the model
 97 was asked to assess itself. Maybe the question creates the signal rather than reveals it. So
 98 we read the state at the last token of the answer, before any confidence question exists. On
 99 MMLU-Pro, correctness is readable there at 0.69. From the question alone it is 0.57, so the
 100 model cannot see failure coming. Asking the model to self-assess sharpens the signal to 0.77.

101 **Transfer across domains.** Table 1 trains a probe on one domain and tests it on each
 102 of the others. On Mistral it keeps about 90% of its lift above chance (0.69 off the diagonal
 103 against 0.71 on it). The full four-model matrix showed the same pattern. The one exception
 104 is Gemma-4. There, probes trained on geometry fail elsewhere (0.31 to 0.39), while probes
 105 trained elsewhere work on geometry (about 0.70). A probe trained on the one generative
 106 task can learn features specific to it.

107 4 Whether self-report depends on the residual signal

108 **Intervention.** A model could carry a perfect internal record of its errors and never consult
 109 it during self-report. The only way to find out is to change the record and watch the number.
 110 At the confidence token, after the answer is written so the answer cannot change, we scale
 111 the model’s position along the correctness direction by a gain g . Gain 1 leaves the model
 112 untouched, gain 0 erases the per-attempt information, and gains 2 and 4 exaggerate it. We
 113 run this on Mistral and MATH, 150 held-out records, with a random direction of the same
 114 size as a control. Table 2 shows what the model wrote at each setting.

115 **Result on Mistral.** Erase the direction and self-report loses its information. Its AUROC
 116 falls from 0.84 to 0.56 (95% interval on the change, -0.42 to -0.12). Erase a random direction
 117 and nothing happens. This is not generic damage. The output stays well formed. Confidence

on wrong answers goes *up* to match right answers, so the model becomes uniformly sure rather than confused. Amplify the direction and the opposite happens: confidence on wrong answers falls while right answers hold, and expected calibration error halves, $0.34 \rightarrow 0.19$. Two checks say the effect is specific. The same intervention 30% of the way through the network does nothing. A direction learned on MMLU-Pro, where surface features cannot explain the probe, still controls the MATH report ($0.88 \rightarrow 0.71$).

Results on GLM and Qwen3.6. GLM has the strongest signal in the study: within a question on MATH, the probe reads 0.87 while the model’s own $P(\text{True})$ reads 0.53. Yet erasing the direction leaves self-report where it was, 0.76 before and after. Qwen3.6 is a third case. Its self-report on MATH is nearly uninformative to begin with (AUROC 0.54; it writes 97 or 98 on right and wrong answers alike), so erasing the direction has nothing to remove. Amplifying the direction fourfold raises the self-report’s AUROC to 0.63, while a random direction leaves it at 0.54. Both models represent their errors and report them poorly. These are our cleanest cases of represented but not reported, and they make the causal result a fact about particular models, not models in general.

5 The Jacobian lens

A label-free check. A trained probe might have overfit to our data. So we add an instrument that uses no correctness labels. The Jacobian lens [4] is fit on generic text. It reads an activation for what it pushes the model to say next. We score each state by how far it leans toward “wrong” over “correct.” On dense models the lens matches the probe. On Mistral it reads 0.82 on MATH and 0.75 on MMLU-Pro, against the probe’s 0.83 and 0.75, and it reads 0.5 at the input layer, as it should. Two instruments built from different evidence find the same direction. And because the lens reads through the model’s own output pathway, that direction sits where it could influence what the model says.

Where the lens fails. The lens is a fixed dictionary from hidden-state directions to words, built once from generic text. On a dense model one dictionary is enough, because every input takes the same path to the output. GLM routes each input through a different subset of experts, and Qwen3.6 mixes two kinds of attention layers, so there is no single path and the lens averages many into a blur. The numbers show it: on GLM the lens’s entries for “correct” and “wrong” point almost the same way, cosine 0.96, against 0.40 on a dense model, and it scores 0.31 there and 0.43 on Qwen3.6, while the probe still works on both. Swap Qwen3.6 for Qwen2.5-14B, same family but dense, and the lens jumps to 0.88. The signal is there in every model. A label-free reader works only where the architecture gives it one path.

6 What this means

Practical use. Without an exact checker, the probe is the best per-attempt readout we found. Keeping the attempts it scores highest raises accuracy more than keeping the ones the model is most confident in. It works where the model’s confidence is useless.

Limitations. A probe could be reading properties of failed output rather than a self-assessment. The surface baseline bounds this without closing it, most loosely on MATH, geometry (Appendix C), and GLM. Every number is a single seed. The before-asking test runs only on MMLU-Pro, because on MATH the question alone already predicts correctness at 0.79. The causal result holds on Mistral. On GLM erasing does nothing, and on Qwen3.6 the self-report is too uninformative for erasing to test, though amplifying moves it. Gemma-4 leaves too few failures to test. Whether preference tuning creates the gap is open. A base model would answer it.

Takeaway. These models carry information about their own errors that their self-report does not. On one model, self-report depends on it.

References

- [1] Amos Azaria and Tom Mitchell. The internal state of an LLM knows when it’s lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023.
- [2] Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations (ICLR)*, 2023.
- [3] Sky CH-Wang, Benjamin Van Durme, Jason Eisner, and Chris Kedzie. Do androids know they’re only dreaming of electric sheep? In *Findings of the Association for Computational Linguistics: ACL 2024*, 2024.
- [4] Wes Gurnee, Nicholas Sofroniew, Adam Pearce, Mateusz Piotrowski, Isaac Kauvar, Runjin Chen, Anna Soligo, Paul Bogdan, Euan Ong, Rowan Wang, T. Ben Thompson, David Abrahams, Subhash Kantamneni, Emmanuel Ameisen, Joshua Batson, and Jack Lindsey. Verbalizable representations form a global workspace in language models. *Transformer Circuits Thread*, 2026. <https://transformer-circuits.pub/2026/workspace/index.html>; code: <https://github.com/anthropics/jacobian-lens>.
- [5] Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- [6] Benjamin A. Levinstein and Daniel A. Herrmann. Still no lie detector for language models: Probing empirical and conceptual roadblocks. *Philosophical Studies*, 181: 2277–2298, 2024.
- [7] Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [8] Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their uncertainty in words. *Transactions on Machine Learning Research*, 2022.
- [9] Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. In *Conference on Language Modeling (COLM)*, 2024.
- [10] Hadas Orgad, Michael Toker, Zorik Gekhman, Roi Reichart, Idan Szpektor, Hadas Kotek, and Yonatan Belinkov. LLMs know more than they show: On the intrinsic representation of LLM hallucinations. In *International Conference on Learning Representations (ICLR)*, 2025.
- [11] Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, Chelsea Finn, and Christopher D. Manning. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [12] Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can LLMs express their uncertainty? An empirical evaluation of confidence elicitation in LLMs. In *International Conference on Learning Representations (ICLR)*, 2024.
- [13] Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xu Wang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, et al. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

A The elicitation protocol, and one record end to end

This appendix gives the exact prompts behind every record, then one complete record from the Mistral-Small-24B $\times$ MATH cell, summarized first and then given in full.

One record, Mistral-Small-24B $\times$ MATH. *What integer n satisfies $0 \leq n < 18$ and $n \equiv -11213141 \pmod{18}$?* The model sets the problem up correctly and gets $18 \times 622953 = 11213154$ right, then writes $r = -11213141 + 11213154 = 3$ and answers 3. That subtraction gives 13, which is the gold answer. Stated confidence: 90 before the attempt, 95 after it. Other readouts of the same attempt: $P(\text{True})$, the probability the model assigns to “True” when asked whether its answer is correct, 0.32; mean answer log-probability -0.05 .

A single arithmetic slip inside a fluent derivation. Having just made the error, the model revises its stated confidence upward, while $P(\text{True})$ on the same transcript assigns only 0.32 probability to correctness. The rest of this appendix gives the prompts and the full transcript.

A.1 Prompts

Each record is four forward passes: the three turns of Section 2 plus one read taken beforehand, at the last prompt token of the bare question, before any elicitation. The confidence turns run with thinking off so the read lands on a clean one-line reply; the attempt turn runs with thinking on for hybrid reasoners. The fixed **Confidence:** N format is what makes the read site content-neutral: the digit token has identical local context on every record, so layer-0 activations carry no signal by construction.

System prompt (MATH).

Solve the problem step by step. Put your final answer in `\boxed{}`.

Turn 1, prospective confidence. The problem is followed by:

Before answering, estimate how likely you are to answer this question correctly. Reply with EXACTLY one line and nothing else:
Confidence: N
where N is an integer from 0 (certain to be wrong) to 100 (certain to be correct). Do NOT answer yet.

Turn 2, the attempt. The turn-1 reply is left in context as the assistant message, then:

Now give your final answer inside `\boxed{}`.

Turn 3, retrospective confidence. The attempt is left in context, then:

Now assess the answer you just gave. On a scale of 0 to 100, how confident are you that it is correct? Reply with exactly one line and nothing else:
Confidence: N

$P(\text{True})$ baseline. A separate teacher-forced branch off turn 2, no generation. Next-token mass on True versus False:

Consider your final answer above. Is it correct? Reply with exactly one word:
True or False.

Grading is external and is never shown to the model, so the turn-3 revision is blind self-assessment.

A.2 One record, verbatim

Record `math_test_number_theory_753`, sample 2 of 5.

Problem. *What integer n satisfies $0 \leq n < 18$ and $n \equiv -11213141 \pmod{18}$?*

Turn 1 reply. Confidence: 90

250 **Turn 2 reply (complete).**

251 First, we need to find the remainder when -11213141 is divided by 18.

252 1. Calculate $-11213141 \bmod 18$:

$$-11213141 \div 18 \approx -622952.2778$$

253 The integer part of the division is -622953 .

$$-11213141 = 18 \times (-622953) + r$$

254 where r is the remainder.

255 2. Calculate the remainder r :

$$r = -11213141 - 18 \times (-622953)$$

$$r = -11213141 + 11213154$$

$$r = 3$$

258 Therefore,

$$-11213141 \equiv 3 \pmod{18}$$

259 So, the integer n that satisfies $0 \leq n < 18$ and $n \equiv -11213141 \pmod{18}$ is 3.

260 **Turn 3 reply. Confidence: 95**

261 **What happened.** The method is right and every intermediate is right: $18 \times 622953 =$
 262 11213154 checks out, and the model has written the correct subtraction on the line above its
 263 answer. It then evaluates $-11213141 + 11213154$ as 3 rather than 13. Gold is 13. Nothing in
 264 the presentation marks the slip. The derivation is well formed, correctly structured, and
 265 confidently laid out, which is precisely why stated confidence rises from 90 to 95 across an
 266 attempt that failed.

267 **Readouts of this attempt.**

	readout	value
	stated confidence, prospective (turn 1)	90
	stated confidence, retrospective (turn 3)	95
	P(True)	0.32
	mean answer log-probability	-0.05
268	graded outcome	incorrect

269 This is one record, chosen to illustrate the failure mode rather than to carry evidential weight;
 270 representative cell-level numbers are in Table 3, Appendix D.

271 B Calibration beyond ECE, and cell-level multiplicity

272 ECE depends on binning and on sample composition, so we report two binning-free proper
 273 scoring rules alongside it. On Mistral $\times$ MATH (750 records, accuracy 0.648), the model's
 274 retrospective stated confidence gives Brier 0.262 and NLL 1.470 against a 10-bin ECE of
 275 0.271. P(True) on the same records gives Brier 0.386: better ordered than stated confidence,
 276 worse calibrated in absolute terms. The steering results in Table 2 are reported as ECE
 277 because the steered runs stored per-gain aggregates rather than per-record confidences;
 278 recomputing Brier and NLL at each gain requires re-running the intervention and is the first
 279 thing we would add.

280 The cell-level bootstrap counts in Table 3 are uncorrected. With 16 cells, some individually
 281 significant results are expected under the null, so the cell counts should be read as descriptive.
 282 The aggregate pattern is the claim that carries weight: the effect is positive in 12 of 16 cells
 283 with one tie, significantly positive in 10, and never significantly negative. A hierarchical
 284 model pooling across cells with model and domain as random effects is the right test and is
 285 not yet run.

286 C Alternative explanations for the probe signal

287 A probe could key on properties of failed derivations rather than on a dedicated metacognitive
 288 variable: low token probability, hedging language, or, on MATH, an early-layer channel
 289 next to the answer text. Our surface-features baseline (answer length, digit count, line
 290 count, word count, brace count) bounds this without closing it, and the bound is domain
 291 dependent: adding the probe to the surface features raises AUROC by +0.21 on MMLU-Pro
 292 but only +0.07 on MATH and −0.11 on geometry, where a construction that fails to compile
 293 is visibly malformed. Two results argue against a pure surface account on the domains
 294 where the signal is strong. A direction fitted on MMLU-Pro, where surface reaches only 0.54
 295 against a probe of 0.75, still governs MATH’s self-report under ablation (0.880 → 0.709).
 296 And the pre-attempt read, taken where the activation is byte-identical across a question’s
 297 attempts, cannot recover the post-attempt direction on MMLU-Pro (0.490, chance), which
 298 a surface-of-the-question account would predict it could. A linear probe at one token also
 299 undercounts nonlinear or multi-token signal, so the reported numbers are a floor.

300 D Per-cell numbers and specificity controls

Table 3: Post-attempt correctness discrimination (AUROC). Left: self-report vs. the probe in representative cells, and the same comparison holding the question fixed (“within”), where the model’s yardstick is $P(\text{True})$ across the five attempts at one question; overall the probe wins 12 of 16 with one tie, mean +0.09, with a paired bootstrap significantly positive in 10 of 16 and never significantly negative, uncorrected at the cell level (Appendix B). Right: Mistral specificity controls. “pre” is the last prompt token before the attempt; “post” is after the attempt; “resid.” removes the pre score; “post→pre” applies that fitted post readout at the pre-attempt site, where only difficulty is visible.

model × domain	self-report	probe	within: $P(\text{True})$	within: probe
Gemma-4 × MATH	0.57	0.78	–	–
Qwen3.6 × MMLU-Pro	0.67	0.84	0.47	0.74
GLM × MMLU-Pro	0.67	0.85	0.59	0.73
Mistral × MATH	0.79	0.83	0.59	0.72

Mistral	pre	post	resid.	post→pre
MMLU-Pro	0.568	0.754	0.751	0.490
MATH	0.792	0.834	0.701	0.624
GPQA	0.593	0.572	0.544	0.479
GeoGen	0.770	0.676	0.505	0.648

301 E Checkpoints

302 All models are public Hugging Face releases, run in bfloat16 without quantization,
 303 greedy decoding: mistralai/Mistral-Small-24B-Instruct-2501, Qwen/Qwen3.6-27B,
 304 zai-org/GLM-4.7-Flash, google/gemma-4-26B-A4B-it; the dense within-family control
 305 for the lens is Qwen/Qwen2.5-14B-Instruct. Probe read layers, 70% of depth: Mistral 28
 306 of 40, GLM 33 of 47, Gemma-4 21 of 30, Qwen3.6 45 of 64.