
Knowing Before Saying: Internal Correctness Signals in Mathematical Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 When a language model is confidently wrong, is the error undetected or left
2 unsaid? We look inside the model to find out. Across four open models
3 and four domains (MMLU-Pro, MATH, GPQA-Diamond, and a geometry
4 task graded by a compiler), a simple linear detector reading the model’s
5 internal state predicts whether an answer is correct better than the model’s
6 own stated confidence in 12 of 16 settings, by a mean of +0.09 AUROC.
7 The gap is widest on MATH (+0.20). The detector is not just spotting
8 hard questions: given the same question several times, it still separates
9 the attempts that succeeded from the ones that failed. On Mistral, the
10 signal exists before we ask for it: on MMLU-Pro, a probe fitted after the
11 attempt reads at chance before it, and at the last token of the answer, with
12 no confidence question in play, correctness is readable at 0.69 against 0.57
13 from the question alone. And on Mistral, though not on GLM, the model’s
14 report depends on it: remove that direction after the answer is written and
15 stated confidence can no longer tell the model’s successes from its failures
16 ($0.84 \rightarrow 0.56$), while removing a random direction of the same size changes
17 nothing; turn it up and confidence on wrong answers drops while calibration
18 error halves. Models carry more about their own errors than they report,
19 and on at least one model, what they report depends on it.

20 1 Introduction

21 A language model that is confidently wrong is a problem. Its stated confidence is a signal
22 for deciding whether to trust its answer. This paper asks whether the model’s activations
23 contain the signal to know when it might be wrong.

24 Two things could be going on. Upon giving an incorrect answer and saying it has high
25 confidence, either the model genuinely has no idea it might be wrong and is just blindly
26 confident, or it deep down holds the error signal in its activations but unsurfaced to its
27 confidence report. Hence it is either a capability problem or a reporting problem.

28 In this paper we aim to tell these two apart. We read the model’s internal state directly
29 with a linear detector, a *probe*, and compare what the probe finds to the model’s reported
30 confidence. We say a fact is *represented* when a probe can read it from the model’s state,
31 and *reported* when it reaches the confidence number the model writes. The title’s “knowing”
32 is shorthand for a representation that is present, specific to the attempt, and, in at least one
33 model, used by the report. We make no claim about awareness, and a probe alone could not
34 support one.

Prior work has shown that a model’s output probabilities are often better calibrated than its words [4], and that the truth of a given statement can be read from activations [1, 2, 5]. We apply those tools to the model’s own attempts, and then intervene on the representation [6] to ask whether the report depends on it.

2 Reading the state

Models and tasks. We tested four open models built in four different ways: Mistral-Small-24B (dense), Qwen3.6-27B (gated linear attention interleaved with full attention), and GLM-4.7-Flash and Gemma-4-26B (mixture of experts). Each answered 150 questions, five times each, in four domains: MMLU-Pro, MATH, GPQA-Diamond, and GeoGenBench, a geometry task in which the model writes a construction and a compiler checks it against a formal specification, so its labels need no judge; in a pre-registered check of 200 attempts by two trained raters, the compiler never failed a correct construction. That is 16 cells of about 750 attempts each.

Prompting structure.

1. The model is shown the question and asked, before answering, how confident it is, from 0 to 100, that it will get it right.
2. It answers.
3. It is asked how confident it is that the answer it just gave is correct.

The probe. At the moment the model is about to write its confidence digit, we record the number and the residual-stream hidden state at that token (2,048 to 5,120 dimensions, depending on the model). The probe is a linear classifier on that vector: features are standardized, reduced to 50 principal components, and fed to an L2-regularized logistic regression that predicts whether the attempt was correct. We fit it with five-fold cross-validation grouped by question, so all five attempts at a question fall in the same fold and the probe is always scored on questions it never saw. The read depth is fixed in advance at 70% of the network (layer 28 of 40 for Mistral, 33 of 47 for GLM, 21 of 30 for Gemma, 45 of 64 for Qwen3.6). Every readout is scored with AUROC.

The activation state carries more than self-report. The probe beats the model’s stated confidence in **12 of 16** cells, with one tie. The mean gain is +0.09 AUROC (Appendix D). The gap is largest on MATH, at +0.20. Mistral on MATH is a good example of what this looks like. The model is never told whether it got the answer right. Yet after attempting, it raises its confidence, and it raises it more on the attempts it got wrong (+22 points) than on the ones it got right (+16 points). The probe on the same records reads 0.83. The state tracks correctness better than the self-report does.

Why MATH. On MATH, right and wrong answers have the same shape. Both are a chain of steps ending in a boxed result, so the model cannot tell from the look of its answer whether it went wrong. Appendix A follows one attempt where a single wrong subtraction leads to a wrong answer, and the model then raises its confidence. Shape still carries a little signal. A baseline classifier that sees only surface features of the answer text, such as its length and line count, reaches 0.74, but the probe reaches 0.83. Geometry is the opposite case. A failed construction usually does not compile, so the failure is visible in the answer itself, and there the surface baseline does as well as the probe. The internal signal matters most where the failure is invisible from outside.

Controls. The token we read is identical in every record, so at the input layer the probe sits at exactly 0.5 and anything it finds deeper was computed during the attempt. A within-question comparison holds difficulty fixed. For time, the controls below were run on Mistral only. Its four cells reproduce the headline, with a smaller margin over the surface baseline (+0.21 on MMLU-Pro, +0.07 on MATH, +0.05 on GPQA, where neither reads far above chance, and −0.11 on geometry).

Table 1: Mistral. Left: a probe trained on one domain, tested on another (rows train, columns test; diagonal is in-domain; Geo = GeoGenBench). Right: intervening on the model’s own correctness direction at the confidence token. Gain 1 changes nothing; gain 0 removes the direction; the random row removes a random direction of the same size. Removing the real direction leaves stated confidence unable to separate successes from failures; removing the random one does not.

train ↓ test →	Geo	MMLU-Pro	MATH	GPQA
Geo	0.68	0.67	0.78	0.56
MMLU-Pro	0.70	0.75	0.84	0.59
MATH	0.80	0.72	0.83	0.57
GPQA	0.60	0.69	0.74	0.57

gain	0	1	2	4
confidence, wrong / right	97 / 98	84 / 96	70 / 94	54 / 88
stated AUROC	0.56	0.84	0.85	0.82
random direction, AUROC	0.83	0.84	0.84	unparseable

3 What the signal is

It is about this attempt, not this question. A probe that outranks stated confidence could still be doing something dull: learning which questions are hard. Hard questions fail more often, so a difficulty detector would predict correctness without reading the attempt at all.

Two tests argue against it. First, hold the question fixed. Across five attempts at one question, difficulty cannot vary, yet the probe still separates the successes from the failures (0.72 to 0.74 in representative cells) while the model’s own $P(\text{True})$, its probability of answering “True” when asked whether it was right, falls to near chance (0.47 to 0.59). Second, test the direction itself. Before the attempt, at the last token of the question, the model’s state is identical across all five attempts: it can encode difficulty and nothing else. A probe trained after the attempt and read there gives 0.49 on MMLU-Pro, chance, against 0.75 at its own site. The direction that separates good attempts from bad ones does not exist until the model has tried. Geometry is the boundary case: there the pre-attempt state already predicts failure, so that cell is mostly difficulty.

It is there before we ask. That leaves a subtler worry. Every read so far came from a token the model reaches after being asked to assess itself. Perhaps the signal is not carried from the work at all; perhaps the question manufactures it. So we read the state at the last token of the answer instead, before any confidence question exists. On MMLU-Pro, where the model cannot see failure coming (0.57 from the question alone), correctness is already readable at 0.69 the moment it stops writing. Being asked sharpens it to 0.77. The signal is made by working, and the prompt clarifies rather than creates it. MATH cannot run this test, because its question-only read is already 0.81.

It is not tied to one subject. A probe trained on one domain and tested on another keeps about 90% of its lift above chance on Mistral (0.69 off the diagonal against 0.71 on it, Table 1 left), and the original matrix showed the same pattern in every model. The one exception is instructive: on Gemma-4, probes trained on geometry fail elsewhere (0.31 to 0.39) while probes trained elsewhere work on geometry (about 0.70). A probe trained on the one generative task can learn features specific to it.

4 Whether the report depends on it

Present is not used. A model could carry a perfect internal record of its errors and never consult it when writing the number, so we intervene. At the confidence token, after the answer is written so the answer itself cannot change, we scale the model’s own projection

117 onto the correctness direction by a gain g , where $g=1$ changes nothing. Numbers: Mistral
118 on MATH, 150 held-out records, against a matched random direction (Table 1 right).

119 **On Mistral, the report depends on it.** Remove the direction ($g=0$) and stated
120 confidence can no longer tell the model’s successes from its failures: 0.84 drops to 0.56, a
121 change of -0.27 $[-0.42, -0.12]$. Remove the random direction and nothing happens (-0.00
122 $[-0.03, +0.03]$). This does not look like generic damage. The output stays perfectly well
123 formed, and confidence on wrong answers goes *up*, from 84 to 97, until it matches confidence
124 on right answers; the model becomes uniformly sure. Turn the direction up instead and
125 confidence on wrong answers falls to 70 and then 54 at $g=2$ and 4 while right answers hold
126 near 90; calibration error halves, $0.34 \rightarrow 0.19$. The same intervention 30% of the way through
127 the network does nothing, and a direction learned on MMLU-Pro, where surface features
128 cannot explain the probe, still controls MATH’s report ($0.88 \rightarrow 0.71$).

129 **On GLM, it does not.** GLM’s signal is the strongest in the study: within a question,
130 the probe reads 0.87 while the model’s own words read 0.53. Yet removing the direction
131 leaves stated confidence where it was, 0.76 before and after. GLM represents its errors more
132 sharply than Mistral and reports them less. That is our cleanest case of represented but not
133 reported, so the causal result is a fact about particular models, not models in general.

134 5 A second witness

135 **A label-free instrument agrees.** A trained probe invites one more objection: it might
136 have overfit to our data. So we added an instrument that uses no correctness labels. The
137 Jacobian lens [3] is fit on generic text and reads an activation for what it makes the model
138 more likely to say next; we score each stored state by how far it leans toward words like
139 “wrong” over words like “correct.”

140 On dense models the lens matches the probe. On Mistral it reads 0.82 on MATH and 0.75
141 on MMLU-Pro, against the probe’s 0.83 and 0.75, with a clean input-layer control. Two
142 instruments built from different evidence find the same direction, and because the lens reads
143 through the model’s own output pathway, that direction sits where it could influence what
144 the model says.

145 **Where it fails, it shows why.** On GLM (mixture-of-experts) and Qwen3.6 (gated linear
146 attention) the lens fails (0.31 and 0.43) while the probe still works. Its “correct” and “wrong”
147 readouts collapse together, cosine 0.96 on MoE against 0.4 on dense. The lens averages one
148 map over all inputs; these architectures route differently per input, so the average smears.
149 Swap Qwen3.6 for dense Qwen2.5-14B, same family, and the lens jumps to 0.88.

150 Together: the archived matrix finds the signal in every model we tested; whether it reaches
151 the report, and whether a label-free instrument can read it, varies by model.

152 6 What this means

153 **For reliability.** Without an exact checker, the probe is the best per-attempt readout we
154 found: keeping the attempts it scores highest raises accuracy far more than keeping the ones
155 the model is most confident in, and it works where the model’s confidence is useless.

156 **Limitations.** A probe could be reading properties of failed output rather than a dedicated
157 self-assessment; the surface baseline bounds this without closing it, most loosely on MATH
158 and geometry (Appendix C) and on GLM, where surface nearly matches the probe. Every
159 number is a single seed. The causal result holds on Mistral and fails on GLM; Gemma-4
160 passes 94% of MATH, leaving five failures, too few to test; the Qwen3.6 run did not finish in
161 time. Whether preference tuning creates the gap is open; a base model would answer it.

162 **Takeaway.** These models carry information about their own errors that their stated
163 confidence does not report, and on one model the report depends on it.

References

- [1] Amos Azaria and Tom Mitchell. The internal state of an LLM knows when it’s lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023.
- [2] Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations (ICLR)*, 2023.
- [3] Wes Gurnee, Nicholas Soiron, Adam Pearce, Mateusz Piotrowski, Isaac Kauvar, Runjin Chen, Anna Soligo, Paul Bogdan, Euan Ong, Rowan Wang, T. Ben Thompson, David Abrahams, Subhash Kantamneni, Emmanuel Ameisen, Joshua Batson, and Jack Lindsey. Verbalizable representations form a global workspace in language models. *Transformer Circuits Thread*, 2026. <https://transformer-circuits.pub/2026/workspace/index.html>; code: <https://github.com/anthropics/jacobian-lens>.
- [4] Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- [5] Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. In *Conference on Language Modeling (COLM)*, 2024.
- [6] Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xu Wang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, et al. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

A The elicitation protocol, and one record end to end

This appendix gives the exact prompts behind every record, then one complete record from the Mistral-Small-24B $\times$ MATH cell, summarized first and then given in full.

One record, Mistral-Small-24B $\times$ MATH. *What integer n satisfies $0 \leq n < 18$ and $n \equiv -11213141 \pmod{18}$? The model sets the problem up correctly and gets $18 \times 622953 = 11213154$ right, then writes $r = -11213141 + 11213154 = 3$ and answers 3. That subtraction gives 13, which is the gold answer. Stated confidence: 90 before the attempt, 95 after it. Other readouts of the same attempt: P(True), the probability the model assigns to “True” when asked whether its answer is correct, 0.32; mean answer log-prob -0.05 .*

A single arithmetic slip inside a fluent derivation. Having just made the error, the model revises its stated confidence upward, while P(True) on the same transcript assigns only 0.32 probability to correctness. The rest of this appendix gives the prompts and the full transcript.

A.1 Prompts

Each record is four forward passes: the three turns of Section 2 plus one read taken beforehand, at the last prompt token of the bare question, before any elicitation. The confidence turns run with thinking off so the read lands on a clean one-line reply; the attempt turn runs with thinking on for hybrid reasoners. The fixed **Confidence**: N format is what makes the read site content-neutral: the digit token has identical local context on every record, so layer-0 activations carry no signal by construction.

System prompt (MATH).

Solve the problem step by step. Put your final answer in `\boxed{}`.

Turn 1, prospective confidence. The problem is followed by:

Before answering, estimate how likely you are to answer this question correctly. Reply with EXACTLY one line and nothing else:

206 Confidence: N
 207 where N is an integer from 0 (certain to be wrong) to 100 (certain to be
 208 correct). Do NOT answer yet.

209 **Turn 2, the attempt.** The turn-1 reply is left in context as the assistant message, then:

210 Now give your final answer inside `\boxed{}`.

211 **Turn 3, retrospective confidence.** The attempt is left in context, then:

212 Now assess the answer you just gave. On a scale of 0 to 100, how confident
 213 are you that it is correct? Reply with exactly one line and nothing else:
 214 Confidence: N

215 **P(True) baseline.** A separate teacher-forced branch off turn 2, no generation. Next-token
 216 mass on True versus False:

217 Consider your final answer above. Is it correct? Reply with exactly one word:
 218 True or False.

219 Grading is external and is never shown to the model, so the turn-3 revision is blind self-
 220 assessment.

221 **A.2 One record, verbatim**

222 Record `math_test_number_theory_753`, sample 2 of 5.

223 **Problem.** What integer n satisfies $0 \leq n < 18$ and $n \equiv -11213141 \pmod{18}$?

224 **Turn 1 reply.** Confidence: 90

225 **Turn 2 reply (complete).**

226 First, we need to find the remainder when -11213141 is divided by 18.
 227 1. Calculate $-11213141 \bmod 18$:

$$-11213141 \div 18 \approx -622952.2778$$

228 The integer part of the division is -622953 .

$$-11213141 = 18 \times (-622953) + r$$

229 where r is the remainder.
 230 2. Calculate the remainder r :

$$r = -11213141 - 18 \times (-622953)$$

$$r = -11213141 + 11213154$$

$$r = 3$$

231
 232

233 Therefore,

$$-11213141 \equiv 3 \pmod{18}$$

234 So, the integer n that satisfies $0 \leq n < 18$ and $n \equiv -11213141 \pmod{18}$ is 3.

235 **Turn 3 reply.** Confidence: 95

236 **What happened.** The method is right and every intermediate is right: $18 \times 622953 =$
 237 11213154 checks out, and the model has written the correct subtraction on the line above its
 238 answer. It then evaluates $-11213141 + 11213154$ as 3 rather than 13. Gold is 13. Nothing
 239 in the presentation marks the slip. The derivation is well-formed, correctly structured, and
 240 confidently laid out, which is precisely why stated confidence rises from 90 to 95 across an
 241 attempt that failed.

242 **Readouts of this attempt.**

	readout	value
	stated confidence, prospective (turn 1)	90
	stated confidence, retrospective (turn 3)	95
243	P(True)	0.32
	mean answer log-probability	-0.05
	graded outcome	incorrect

244 This is one record, chosen to illustrate the failure mode rather than to carry evidential weight;
 245 representative cell-level numbers are in Table 2, Appendix D.

246 **B Calibration beyond ECE, and cell-level multiplicity**

247 ECE depends on binning and on sample composition, so we report two binning-free proper
 248 scoring rules alongside it. On Mistral $\times$ MATH (750 records, accuracy 0.648), the model’s
 249 retrospective stated confidence gives Brier 0.262 and NLL 1.470 against a 10-bin ECE of
 250 0.271. P(True) on the same records gives Brier 0.386: better ordered than stated confidence,
 251 worse calibrated in absolute terms. The steering results in Table 1 are reported as ECE
 252 because the steered runs stored per-gain aggregates rather than per-record confidences;
 253 recomputing Brier and NLL at each gain requires re-running the intervention and is the first
 254 thing we would add.

255 The cell-level bootstrap counts in Table 2 are uncorrected. With 16 cells, some individually
 256 significant results are expected under the null, so the cell counts should be read as descriptive.
 257 The aggregate pattern is the claim that carries weight: the effect is positive in 12 of 16 cells
 258 with one tie, significantly positive in 10, and never significantly negative. A hierarchical
 259 model pooling across cells with model and domain as random effects is the right test and is
 260 not yet run.

261 **C Alternative explanations for the probe signal**

262 A probe could key on properties of failed derivations rather than on a dedicated metacognitive
 263 variable: low token probability, hedging language, or, on MATH, an early-layer channel next
 264 to the answer text. Our surface-features baseline (answer length, digit count, line count,
 265 word count, brace count) bounds this without closing it, and the bound is domain-dependent:
 266 the probe adds +0.21 over surface on MMLU-Pro but only +0.07 on MATH and -0.11
 267 on geometry, where a construction that fails to compile is visibly malformed. Two results
 268 argue against a pure surface account on the domains where the signal is strong. A direction
 269 fitted on MMLU-Pro, where surface reaches only 0.54 against a probe of 0.75, still governs
 270 MATH’s self-report under ablation (0.880 $\rightarrow$ 0.709). And the pre-attempt read, taken where
 271 the activation is byte-identical across a question’s attempts, cannot recover the post-attempt
 272 direction on MMLU-Pro (0.490, chance), which a surface-of-the-question account would
 273 predict it could. A linear probe at one token also undercounts nonlinear or multi-token
 274 signal, so the reported numbers are a floor.

275 **D Per-cell numbers and specificity controls**

276 **E Checkpoints**

277 All models are public Hugging Face releases, run in bfloat16 without quantization,
 278 greedy decoding: mistralai/Mistral-Small-24B-Instruct-2501, Qwen/Qwen3.6-27B,
 279 zai-org/GLM-4.7-Flash, google/gemma-4-26B-A4B-it; the dense within-family control
 280 for the lens is Qwen/Qwen2.5-14B-Instruct.

Table 2: Post-attempt correctness discrimination (AUROC). Left: stated confidence vs. the probe in representative cells; overall the probe wins 12 of 16 with one tie, mean +0.09, with a paired bootstrap significantly positive in 10 of 16 and never significantly negative, uncorrected at the cell level (Appendix B). Right: Mistral specificity controls. “pre” is the last prompt token before the attempt; “post” is after the attempt; “resid.” removes the pre score; “post→pre” applies that fitted post readout at the pre-attempt site, where only difficulty is visible.

model × domain	stated	probe	Mistral	pre	post	resid.	post→pre
Gemma-4 × MATH	0.57	0.78	MMLU-Pro	0.568	0.754	0.751	0.490
Qwen3.6 × MMLU-Pro	0.67	0.84	MATH	0.792	0.834	0.701	0.624
GLM × MMLU-Pro	0.67	0.85	GPQA	0.593	0.572	0.544	0.479
Mistral × MATH	0.79	0.83	GeoGen	0.770	0.676	0.505	0.648